

## Notes, 7(d)

ECE 606

### Floyd-Warshall

A final example we discuss in Lecture 7 on dynamic programming is an ingenious algorithm for the all-source shortest distances/paths problem.

Previously, we considered single-source shortest distances/paths only. That is, part of our input is a source-vertex  $s$ , and we compute shortest distances and paths to every other vertex from  $s$  only. What if we seek shortest distances and paths between every pair of vertices?


That is, our input is a graph undirected or directed. Suppose the input graph has  $n$  vertices. Assume the vertices are named  $\{1, 2, \dots, n\}$ . Then, our output is two  $n \times n$  matrices,  $D$  and  $\Pi$ .  $D[i, j]$  is the shortest-distance from vertex  $i$  to  $j$ ; it is the special symbol  $\infty$  if no path  $i \rightsquigarrow j$  exists and 0 if  $i = j$ .  $\Pi[i, j]$  is a predecessor of  $j$  in a shortest-path  $i \rightsquigarrow j$ , i.e., of weight  $D[i, j]$ ; it is the special mnemonic NIL if no path  $i \rightsquigarrow j$  exists, or  $i = j$ .

We could simply run Bellman-Ford with every vertex as the source-vertex. As each such run has time-efficiency  $\Theta(VE)$ , we would have a time-efficiency of  $\Theta(V^2E)$ , which is  $\Theta(V^4)$  in the worst-case.

If we know that we have non-negative edges only, then we could run Dijkstra repeatedly with each vertex as the source-vertex with, for example, the priority queue implemented using a binary heap. Then, we would have a worst-case time-efficiency of  $\Theta(VE \lg V) = \Theta(V^3 \lg V)$ .

Floyd-Warshall exploits the optimal substructure of shortest paths the following way. Suppose our vertices are named  $\{1, 2, \dots, n\}$ , i.e.,  $|V| = n$ . Define  $d_{u,v}^{(k)}$  as the shortest distance from  $u$  to  $v$  such that paths are allowed to traverse the vertices  $\{1, 2, \dots, k\}$  only. By “traverse,” we mean, use as intermediate vertices.

For example, for every  $u, v \in V, u \neq v$ ,  $d_{u,v}^{(0)} = w(u, v)$  if  $\langle u, v \rangle \in E$  and  $\infty$  otherwise. For more examples, consider the following graph.



$d_{1,4}^{(0)} = d_{1,4}^{(1)} = \infty$ . But  $d_{1,4}^{(2)} = 4$  on account of the path  $1 \rightarrow 2 \rightarrow 4$ . But that is not a shortest-path in the graph from 1 to 4. It turns out  $d_{1,4}^{(5)} = 2$  on account of  $1 \rightarrow 5 \rightarrow 4$ .

We can write a recurrence for  $d_{i,j}^{(k)}$ , which then immediately suggests a bottom-up algorithm based on dynamic programming.

First, we specify  $w_{i,j}$  for every  $i, j \in V$ , as a generalization of  $w: E \rightarrow \mathbb{R}$ . It is a function,  $w_{i,j}: V \times V \rightarrow \mathbb{R} \cup \{\infty\}$ , where:

$$w_{i,j} = \begin{cases} 0 & \text{if } i = j \\ w(i, j) & \text{if } \langle i, j \rangle \in E \\ \infty & \text{otherwise} \end{cases}$$

Now, our recurrence for  $d_{i,j}^{(k)}$  is as follows:

$$d_{i,j}^{(k)} = \begin{cases} w_{i,j} & \text{if } k = 0 \\ \min\{d_{i,j}^{(k-1)}, d_{i,k}^{(k-1)} + d_{k,j}^{(k-1)}\} & \text{otherwise} \end{cases}$$

We can think of the above recurrence as exploiting the fact that  $d_{i,j}^{(k)} = d_{i,j}^{(k-1)}$  whenever  $i = k$  or  $j = k$ . Because, if we have no negative edge-weight cycles, then every shortest path is simple.

Our corresponding recurrence for  $\pi_{i,j}^{(k)}$ , a predecessor vertex of  $i$  on a path that has weight  $d_{i,j}^{(k)}$ , is as follows:

$$\pi_{i,j}^{(k)} = \begin{cases} \text{NIL} & \text{if } k = 0 \text{ and } w_{i,j} \in \{0, \infty\} \\ i & \text{if } k = 0 \text{ and } w_{i,j} \notin \{0, \infty\} \\ \pi_{i,j}^{(k-1)} & \text{if } k > 0 \text{ and } d_{\mathbf{i}, \mathbf{j}}^{(k-1)} \leq d_{i,k}^{(k-1)} + d_{k,j}^{(k-1)} \\ \pi_{k,j}^{(k-1)} & \text{otherwise} \end{cases}$$

It should now be easy to write down a bottom-up algorithm that uses dynamic programming. We would initialize  $D^{(0)} = [w_{i,j}]$ , and then iterate for all  $k$  from 1 to  $n$ , for all  $i$  from 1 to  $n$ , and for all  $j$  from 1 to  $n$  to progressively construct each  $D^{(1)}, D^{(2)}, \dots, D^{(n)}$ , and simultaneously,  $\Pi^{(k)}$ .

At the moment we compute  $D^{(k)}$ , we need the immediately prior  $D^{(k-1)}$  only, and not, for example,  $D^{(k-2)}, \dots, D^{(0)}$ . So at any moment, we need to maintain two  $n \times n$  matrices only, that correspond to  $D$ .

This algorithm's time-efficiency is  $\Theta(V^3)$ , which is better than both Dijkstra and Bellman-Ford, in the worst-case. Also, it can handle negative edge weights.